

Senior Data Platform Engineer — Platform Architecture | Shanghai, China | Open to Global Remote |
jromeguan@gmail.com | <https://linkedin.com/in/jrome-guan-47888241a> | <https://github.com/jlty258/portfolio>

Guan Chen

Summary

Platform architect with 13+ years designing enterprise data infrastructure. As Technical Architect, scaled Flink platforms 100x, reduced PB-scale batch runtime 40%, and defined privacy computing execution architecture with multi-source data access across 10+ heterogeneous backends.

Experience

Senior Data Platform Engineer

Shanghai Puxin Future Internet Research Institute | Dec 2024 – Present

MIRA / ChainWeaver — Privacy Computing Platform

- Designed storage abstraction and Data Service: 10+ data source types, billion-scale table reads, S3 object storage access
- Defined platform architecture: PQL → Planner → Execution DAG → Scheduler → Execution Engine → Data Service → Storage
- Established execution planning as platform capability, replacing per-workload orchestration with declarative query submission
- Standardized interfaces across planner, scheduler, execution engine, and data service before multi-team integration

Senior Data Platform Engineer / Technical Architect

Shanghai Jiayu Intelligent Robotics | Jun 2020 – Sep 2024

Connected Vehicle Telematics — National-Scale Streaming Platform (China)

- Scaled Flink platform from 10,000 to 1,000,000+ vehicles with stable production output; 10B+ events/day on 1000+ vCore cluster
- Defined layer boundaries: Kafka ingestion, Flink computation, Doris storage, Business API serving — each scaling independently
- Established metadata and data governance as platform primitives across national-scale telemetry datasets
- Enabled operational scalability via Dinky for Flink job lifecycle management

Senior Data Platform Engineer

Ping An Securities | May 2017 – Apr 2020

Enterprise Financial Analytics — Metric Platform (China)

- Reduced 1000+ vCore batch computation from 2.5 hours to 1.5 hours on PB-scale Hadoop storage
- Defined metric platform architecture serving 100,000+ searchable Elasticsearch fields from PB-scale Hive computation

- Solved 2000+ Hive column to 100,000+ searchable field mapping through Map-based dynamic schema expansion
- Established reusable compute-serve pattern adopted across metric onboarding

Data Platform Engineer — Shanghai Jiayin Data Technology

Restaurant intelligence analytics; Storm, Hive, Hadoop

Research Engineer — Shanghai Ship and Shipping Research Institute

Hydrodynamic simulation on distributed computing infrastructure

Skills

Architecture: Platform Design, Execution Planning, Storage Abstraction, Distributed Systems, Metadata, Data Governance

Systems: Kafka, Flink, Spark, Hive, Hadoop, Doris, Hudi, Elasticsearch, HDFS, S3 | **Infrastructure:** Linux, Docker, Kubernetes | **Languages:** Java, Python, C

Education

Master, Naval Architecture and Ocean Engineering — Shanghai Ship and Shipping Research Institute |
Bachelor, Hydraulic Engineering — Tianjin University